

INDICIUM.AI · COMMAND CENTRE

Prove Your Claude Investment Is Paying Off.

Your teams have gone all in on Claude. The Command Centre shows you what that investment is actually delivering: who has adopted it, the real use cases it is running across your teams, and the value they give back, with the risks flagged along the way. It is self-hosted, so conversation content never leaves your estate.

[Take the Tour →](#)

01 · EVERY CLAUDE SURFACE

Claude Is Everywhere. Now You Manage All of It.

Your people use Claude in three ways, Chat, Cowork and Code. The Command Centre brings them together, so you can see how your whole organisation uses AI in one place.

CHAT

Everyday questions to Claude

The day-to-day back-and-forth in the Claude app. See what your people are asking for and which teams lean on it most, without storing Claude's replies.

You see: what's being asked

COWORK

Claude doing the whole job

Cowork is where Claude works like a teammate: it uses skills and connects to your own tools to carry out a whole job across many steps, not just answer a question. You see every step it took and every file it touched.

How deep the work goes

CODE

Engineers building with Claude

Your developers writing software with Claude Code. See what it costs, how much it is shipping, and where. It is measured in real commits, never in the code itself.

What it costs vs. ships

02 · THE GUIDED TOUR

See How Command Centre Works.

A guided click-through of Command Centre, area by area, showing you what each screen tells you and why it matters.

[01 Dashboard](#)[02 Audit Log](#)[03 Alerts](#)[04 Reports](#)[05 S&C Board](#)[06 Cowork Pulse](#)[07 Code Pulse](#)[08 FinOps](#)[09 Use Cases](#)[10 Admin](#)[01 · DASHBOARD](#)

The big picture, at a glance.

A single overview of your whole Claude deployment: how many people are using it, whether adoption is growing, the value it is producing, and anything that needs attention today, so you can read the state of play in under a minute.

- ☒ Adoption, value and spend, all live
- ☒ Tap any number to see the detail behind it
- ☒ Shows what AI wrote and what was measured

The screenshot displays the Indicium AI Command Centre interface. At the top, a navigation bar includes a search bar with the text "app.indicium.ai/" and a "LIVE" status indicator. The left sidebar lists various sections: Dashboard, Audit Log, Alerts (with a count of 26), Reports, S&C Board, Cowork Pulse (with a count of 10), Code Pulse, Use Cases, Overview, and Use Case Governance. The main content area is titled "Dashboard" and "Overview of Claude AI usage across your organisation". It features an "Intelligence Briefing" section, which is highlighted by a blue border and a callout box. The briefing, titled "A plain-English summary, written for you", provides a summary of AI usage: "Adoption is accelerating across the organisation: 12 users active this week (up from 7), with 57 Cowork commits across 4 repositories. Work Claude judged to be mostly feature development and bug fixes, with tasks — 129 sessions involved meaningful workflow steps, and 43 new end-to-end processes have been reused and costs low at £5.50 for the week. However, 36 open alerts warrant attention: 7 high-severity in Cowork and repeated tool failures in the payments service, plus several medium-priority anomalies around PII in prompts. These are early-stage governance gaps rather than systemic breaches, but they need resolution." Below the briefing, there are "AI Recommendations" and "Adoption Momentum" sections. The "Adoption Momentum" section shows "Active users grew 71% week-over-week to 12, with 27 Code commits and 43 confirmed new workflows". The "AI Recommendations" section lists several high-severity alerts, such as "Review and resolve credential exposure alert" and "Investigate payments". At the bottom of the dashboard, there are navigation buttons: "PREV", "PLAY", "NEXT", "Step 1 of 39", and "RESTART".

03 · CAPABILITIES

An Answer for Every Person Who Asks.

The CIO wants adoption. The CFO wants ROI. The CISO wants risk closed. Engineering wants to know it is shipping. Eight purpose-built suites answer each of them from the same trusted data, and every number links to the page that proves it.

[/](#)

Dashboard

A single overview of your whole Claude deployment

[/audit](#)

Audit Log

A complete, searchable record of what was said to Claude

[/alerts](#)

Alerts

Every anomaly, policy breach and risk signal becomes a card that tells you what happened, why it

[/reports](#)

Reports

A dated, defensible summary of AI usage and governance for any period, ready to file

[/library](#)

S&C Board

Every skill and connector your teams have actually

[/pulse](#)

Cowork Pulse

Adoption analytics for

[/pulse/code](#)

Code Pulse

Engineering analytics for Claude Code

[/finops](#)

FinOps

AI spend broken down by cost centre, model

built with AI,
found
automatically
and put in

agent-mode
work

and usage
type, every
month

[/use-cases](#)

Use Cases

The specific
jobs your
teams use AI
for, such as
contract
review or
incident triage

[/admin](#)

Admin

All the
governance
controls
behind one
menu

04 · COWORK PULSE

Is Cowork actually changing how work gets done?


Cowork is Claude's agent mode: hand it a whole task, such as a report, a data pull or a first-pass review, and it plans, works across your files and apps, and returns a finished deliverable. Cowork Pulse measures whether that is genuinely taking hold, looking past how many sessions ran to whether the work is getting more valuable, more delegated, and more resilient.

1 More valuable work, not just more of it

Two signals track whether people are handing Cowork substantial, multi-step work over time, and whether the way they work is genuinely changing, rather than simply running it more often.

2 How much your teams are actually delegating

Shows how much real work is being handed to Cowork and how



far it goes, from quick one-offs to whole tasks run start to finish, with a monthly tier mix that tells you whether the organisation is moving up the curve.

3 The capability you are building, not just the usage

Tracks the repeatable jobs Cowork now runs end to end and how widely they are shared across your teams, so you can see durable capability forming and whether it is spread safely or resting on a few individuals.


See it in the tour →

Measured, not guessed.

Engineering leaders get asked two questions about AI coding tools: what is it costing us, and is it actually helping? Code Pulse answers both from real usage: cost, delegation, shipping and reach, each as a plain-English verdict rather than a wall of charts.

The difference is discipline. Nothing here is inflated: the one forward-looking figure is labelled as a projection, usage that cannot yet be priced is left blank instead of counted as zero, and every card shows the window it was measured over. These are numbers you can put in front of a board without a caveat.

See it in the tour →



06 · FINOPS

Know what Claude costs, to the cost centre.

FinOps breaks your Claude spend down by cost centre, model and usage type, every month. Crucially, it is honest about what is measured and what is estimated, reconciles against your Anthropic invoice, and flags the spend it cannot yet attribute, so the number you give finance is one you can stand behind.


1 Attributed to the team, not just the total

Spend is split by cost centre, model and usage type, with top spenders surfaced, so cost lands with the people driving it rather than a single mystery line item.

2 Honest about estimated versus actual

Every figure is labelled measured or estimated, cache savings are shown in full, and a plain-English note explains exactly why this differs from your Anthropic invoice.

3 A month-end number before month-end



A run-rate projection extrapolates month-to-date spend to a forecast total, and unallocated cost or missing rates are flagged so nothing slips through unbudgeted.

See it in the tour

07 · USE CASES


Find, Track and Even Ideate Your Use Cases Across Your Org.

A use case is a specific task your teams do with AI: reviewing contracts, triaging incidents, drafting proposals. Command Centre finds them in your own usage, including the ones nobody declared, and turns an ideation session or a discovery-call transcript into use cases you can track.

FIND · TRACK · IDEATE

Every use case, in one register.


- 1 Find what AI is already doing**
Surface use cases running across your teams, including the ones a team built on their own and never told anyone about.
- 2 Track each one over time**
Every use case gets an owner and a status that updates itself from real usage, so you can see at a glance which are taking hold and which need a nudge.
- 3 Ideate where AI could do more**
Brainstorm new use cases grounded only in your own usage, never a generic list, to surface the value you are leaving on the table.



★ STRATEGIC USE CASES

Track the ones that matter, by department and org.

Flag your priority use cases as strategic and they roll up into one board-ready view: which are adopted, piloting or at-risk, per team and across the org, with the capacity they have given back.



■ measured: sessions & users

◆ estimated: capacity recovered

[See it in the tour →](#)


08 · SKILLS & CONNECTORS

Most teams have no idea what they have already built.

Skills and connectors get built inside individual accounts and never shared, so the same capability is rebuilt team after team. The Skills & Connectors Board finds every one that is actually in use, gets it claimed and documented, and turns reusable AI capability from something invisible into a shared, ranked asset.

1 Discovered for you, not catalogued by hand

Every skill and connector actually used in production is found automatically, with real usage stats,



including the ones nobody
declared.

2 An Org Capability Score you can move

One number for how much of your
in-production capability is claimed
and documented, so ownership is a
metric you can act on rather than a
mystery.

3 A leaderboard that surfaces your champions

All-time skill and connector use per
person, with medals, so the people
quietly building your AI capability
become easy to spot and back.

See it in the tour

09 · GOVERN WITH CONFIDENCE

Adopt Boldly. The Guardrails Are Already On.

The fastest way to say yes to AI is knowing you can see all of it. Every prompt is scanned the moment it is sent, not overnight and not in a sample. The Command Centre reads intent, catches leaked data against the rules you set, and holds your regulatory position as dated fact. It is all self-hosted, so nothing you are protecting ever leaves your estate.

TWELVE BUILT-IN INTENT CATEGORIES, SCANNED AT INGEST. PLUS YOUR OWN

Access Escalation

Policy Evasion

Destructive Action

Financial Fraud

Data Exfiltration

Covert Exfiltration

Credential Harvesting

Lateral Movement

Shadow IT

Social Engineering

Data Tampering

Reconnaissance

+ your custom categories

critical

high

medium

BEHAVIOUR & INTENT

Read intent, not just keywords

Claude judges what a prompt is trying to do, whether that is exfiltrating data, escalating access, or moving laterally. It does so the instant the prompt is sent, before anyone reviews. Turn categories on or off, add your own, and the server enforces every toggle.

SENSITIVE DATA & SECRETS

Leaks caught inline

17 built-in patterns flag customer PII and live credentials the moment they appear in a prompt. A further 17 always-on file rules, covering keys, certs, .env files, and password dumps, trip the instant a risky file is attached to a session.

REGULATORY POSTURE

GDPR, EU AI Act & ISO 27001

Your DPO, DPIA, the EU AI Act Article 50 notice, and ISO 27001:2022 control ownership are tracked as dated declarations on one completeness meter. Audit readiness becomes a status you can show, not a scramble.

WHAT MAKES THIS DIFFERENT

It uses Claude to police Claude, and it lets you audit the judge itself.

Other tools give you dashboards of what happened. This one reasons about intent in real time and scores your whole organisation's risk in one number. Uniquely, it also exposes the accuracy of the AI that does the grading, so the numbers you report to the board are ones you can defend.

Others govern the pipes. We govern the value.

Databricks and LiteLLM secure the infrastructure Claude runs on. Neither can tell you whether your teams are getting value from it. Command Centre matches them on governance, then adds the adoption, value and use-case layer they simply do not have. Amber notes mark exactly what is partial, so the gap is checkable, not just asserted.

	Command Centre GOVERNANCE & ADOPTION PLATFORM	Databricks MODEL SERVING & GOVERNANCE	LiteLLM GATEWAY / PROXY
--	--	--	----------------------------

GOVERNANCE & CONTROL

Table stakes for a regulated estate. We match the infrastructure players, then go further.

Multi-surface capture (Chat, Code, Cowork)	✓	~ Code + Cowork confirmed, Chat unconfirmed	~ Routed traffic only
Full audit log & session search	✓	~ Delta tables, no search UI	~ Raw logs if enabled, no session view
Native PII & secrets detection	✓	✓	~ Opt-in Presidio bolt-on

	Command Centre GOVERNANCE & ADOPTION PLATFORM	Databricks MODEL SERVING & GOVERNANCE	LiteLLM GATEWAY / PROXY
Behavioural intent taxonomy (12 threat categories)	✓	✗	✗
Alerts & anomaly workflow	✓	 Rate limits & policies, no alert queue	✗
Compliance reports (GDPR, EU AI Act, ISO)	✓	✗	✗
AI-classifier auditability ("watch the watcher")	✓	✗	✗
Enterprise IAM (SSO, SCIM, audit, revoke)	✓	✓	 Admin-UI SSO only, not Claude access
Self-hosted, data never leaves your estate	✓	 True only for Databricks-served traffic	 Self-hostable, but proxy-only visibility

ADOPTION & VALUE INTELLIGENCE

The layer that proves your investment is working. No serving platform or proxy has it, at all.

Executive briefing & org risk score, written by Claude	✓	✗	✗
Delegation depth: are teams handing over real work?	✓	✗	✗
Capacity recovered, estimated in hours and cost	✓	✗	✗
Resilience: would the capability survive a key leaver?	✓	✗	✗

	Command Centre GOVERNANCE & ADOPTION PLATFORM	Databricks MODEL SERVING & GOVERNANCE	LiteLLM GATEWAY / PROXY
Cost per outcome, not just token spend	✓	×	×
Shipping attribution: commits & merged PRs from Claude	✓	×	×
ROI figures labelled measured or estimate, never faked	✓	×	×
Skills & connectors discovery + champions leaderboard	✓	×	×

USE-CASE STRATEGY

Turn scattered AI activity into a tracked portfolio of business outcomes. Unique to Command Centre.

Auto-discovery of live use cases, including undeclared ones	✓	×	×
Strategic tracking per team and across the organisation	✓	×	×
AI-assisted ideation from your own usage and transcripts	✓	×	×

IN ONE LINE

Governance, adoption value and use-case strategy, one Claude-specific stack.

Strong infrastructure governance, no Claude-usage layer.

Cost and routing plumbing, no governance layer.

✓ Native, shipped ~ Partial, see note × Does not exist

3

surfaces, one picture.
Chat, Code & Cowork

100%

self-hosted. Your
data never leaves
your estate

17

sensitive-data
patterns flagged
automatically

SOC2

even the AI that
grades your usage is
auditable

READY WHEN YOU ARE

Stop Guessing What Claude Is Doing for You.

See the Command Centre running on your own data. Delivered by
Indicium.AI, the team that puts enterprise AI into production.

[Replay the Tour →](#)

COMMAND CENTRE

Enterprise AI, Delivered · © 2026 Indicium AI